

Chapter 9: Virtual Memory



Narzu Tarannum

CSE, BU






Intro

- **Virtual memory** – separation of user logical memory from physical memory.
 - Only part of the program needs to be in memory for execution.
 - Logical address space can therefore be much larger than physical address space.
 - Allows address spaces to be shared by several processes.
 - Allows for more efficient process creation.
- Virtual memory can be implemented via:
 - Demand paging
 - Demand segmentation
- **Virtual address space** – logical view of how process is stored in memory






Virtual Memory That is Larger Than Physical Memory





Virtual-address Space

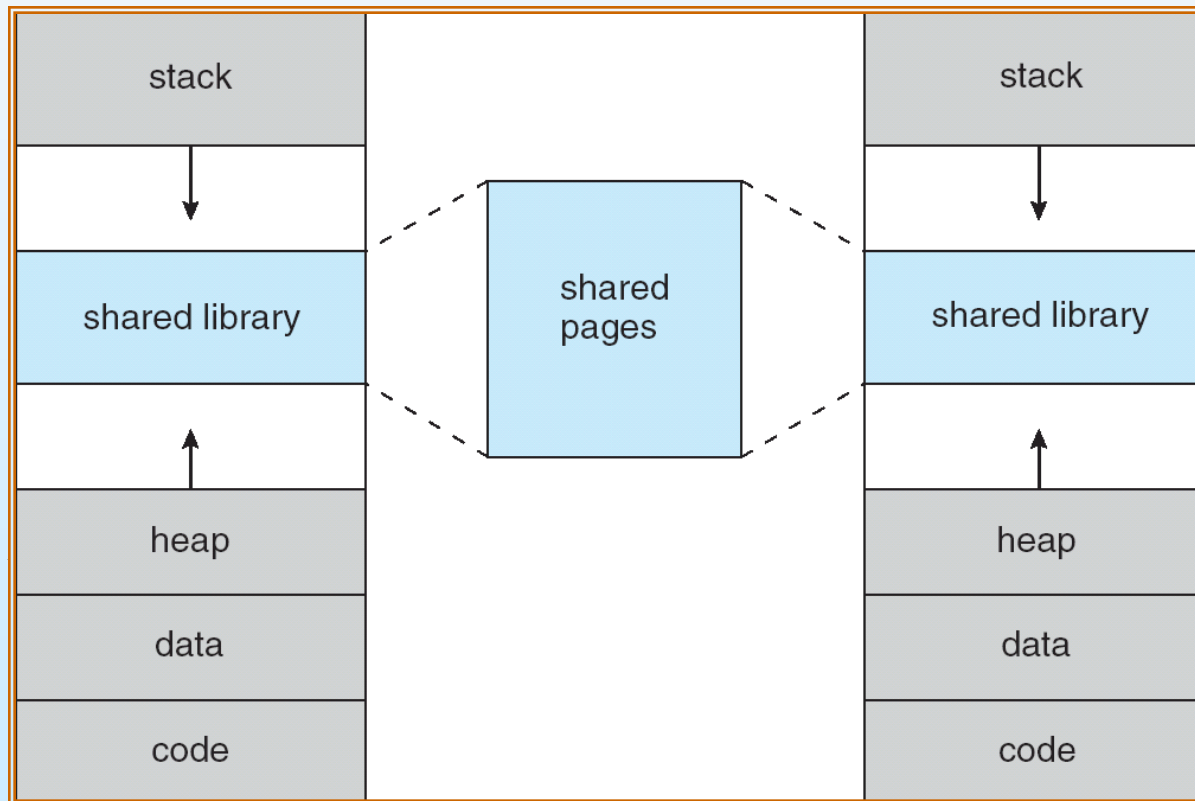
- After compilation, process / executable code presented by its virtual address space
- It is a contiguous address space from address 0 to Max





Shared Library Using Virtual Memory

- Virtual memory also allows files and memory to be shared by two or more processes.





Demand Paging

- Bring a page into memory only when it is needed
- Demand paging follows that pages should only be brought into memory if the executing process demands them. This is often referred to as **lazy evaluation** or **lazy swapper** as only those pages demanded by the process are swapped from secondary storage to main memory.
- Contrast this to **pure swapping/ pure demand paging**, where all memory for a process is swapped from secondary storage to main memory during the process startup.
- Commonly, to achieve this process a page table implementation is used. The page table maps logical memory to physical memory.

Swapper that deals with pages is a **pager**





Demand Paging

■ Advantages

- Less memory needed
- Faster response
- More users/ Degree of multiprogramming increase.
- Less I/O needed
- Reduce memory requirement
- Swap time is also reduce

■ Disadvantages

- Page fault interrupt

■ **locality of reference**, also known as the principle of **locality**, is a term for the phenomenon in which the same values, or related storage locations, are frequently accessed, depending on the memory access pattern.

- Results in reasonable performance from demand paging.

■ Hardware support needed for demand paging

- Page table with valid / invalid bit
- Secondary memory





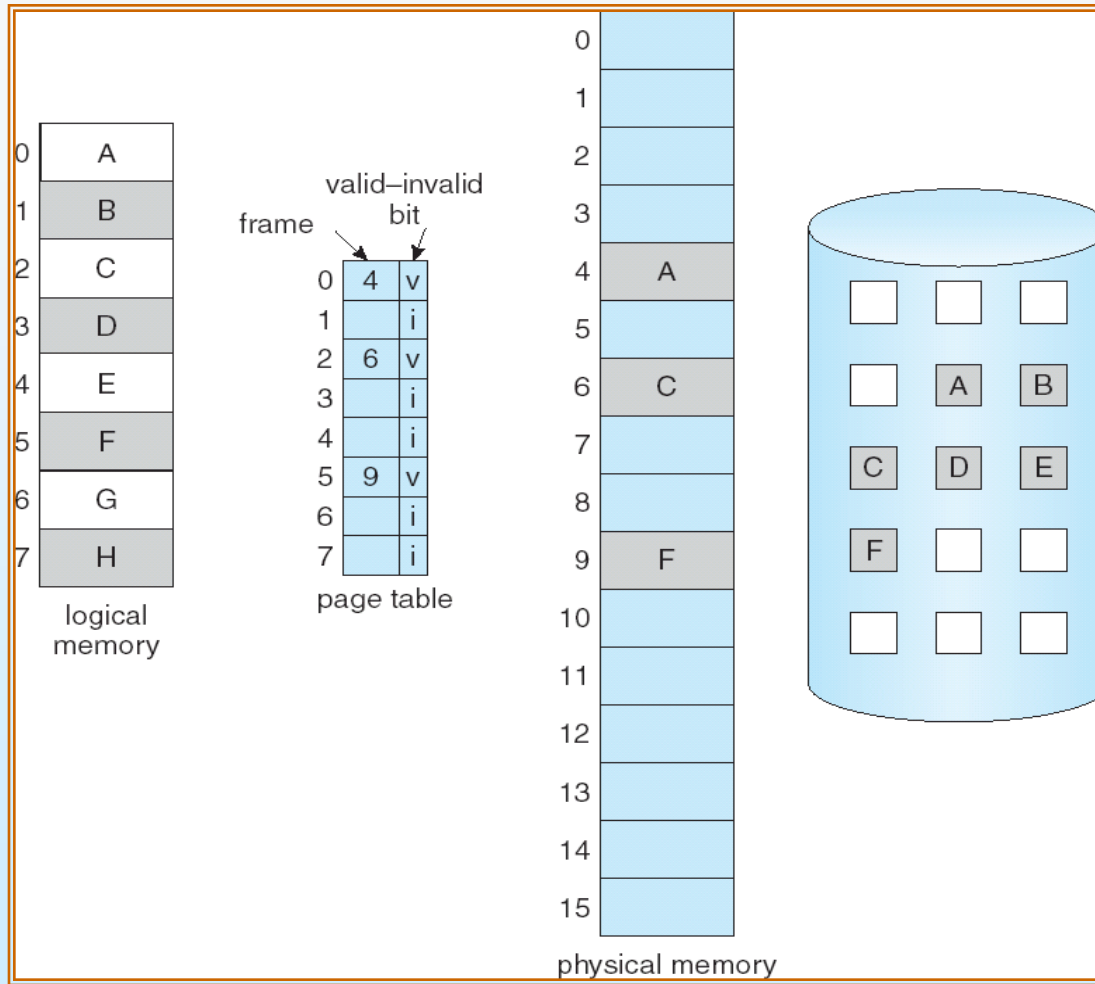
Valid Invalid bit

- The page table uses a bitwise operator to mark if a page is valid or invalid. A valid page is one that currently resides in main memory. An invalid page is one that currently resides in secondary memory. When a process tries to access a page, the following steps are generally followed:
 - Attempt to access page.
 - If page is valid (in memory) then continue processing instruction as normal.
 - If page is invalid then a **page-fault trap / page-fault interrupt** occurs.
 - Page is needed \Rightarrow reference to it
 - invalid reference \Rightarrow abort
 - not-in-memory \Rightarrow bring to memory
 - Restart the instruction that was interrupted by the operating system trap.
- With each page table entry a valid–invalid bit is associated (1 \Rightarrow in-memory, 0 \Rightarrow not-in-memory)





Page Table When Some Pages Are Not in Main Memory






Page Fault

- If there is ever a reference to a page, first reference will trap to OS \Rightarrow **page fault**
- OS looks at another table to decide:
 - Invalid reference \Rightarrow abort.
 - Just not in memory.
- Find empty frame.
- Load page from disk into frame.
- Reset tables, validation bit = 1.
- Restart instruction that caused page fault



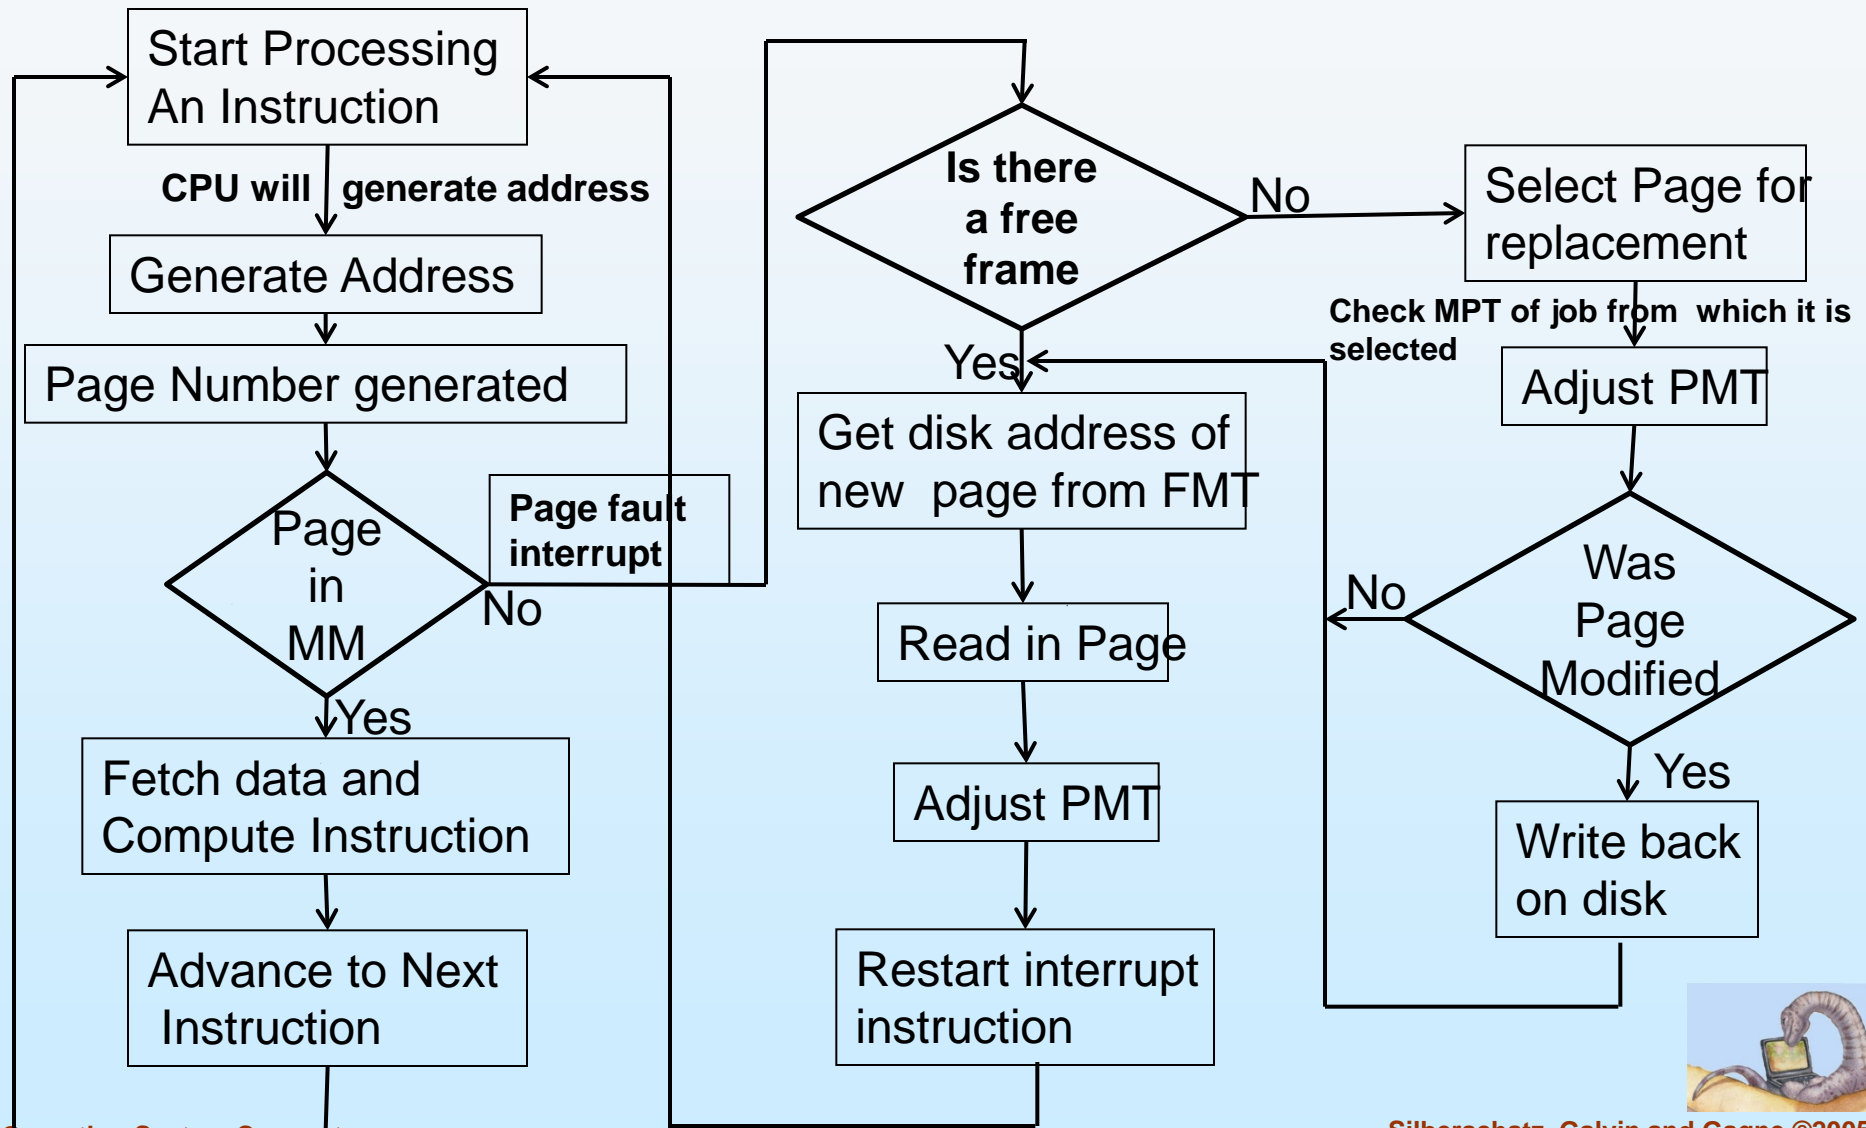


Steps in Handling a Page Fault





Flowchart of Demand Paging






What happens if there is no free frame?

- **Page replacement** – find some page in memory, but not really in use, swap it out
 - algorithm
 - performance – want an algorithm which will result in minimum number of page faults
- Same page may be brought into memory several times





Need For Page Replacement





Performance of Demand Paging

- Demand paging can significantly affect the performance of computer system
- Let's compute the effective access time (EAT) for a demand page memory:
 - If “ p ” be the probability of a page fault where $0 \leq p \leq 1$.
 - ▶ $p = 0$ no page faults
 - ▶ $p = 1$, every reference is a fault
 - We should expect p to be close to 0 or only a few page fault





Performance of Demand Paging

- Then,
 - $EAT = (1 - p) \times ma + p \times \text{page fault time}$
- Let,
 - $ma = \text{memory access time (200 nano-seconds)}$
 - $\text{Average page fault service time} = 8 \text{ mili-seconds}$
* $1\text{ms}=1000000\text{ns}$
- $EAT = (1 - p) \times 200 + p (8 \text{ milli-seconds})$
 $= 200 - 200p + 8,000,000p \text{ (nano-seconds)}$
 $= 200 + 7,999,800p$
- We can see EAT is directly proportional to page fault rate
 - For Example if one access out of one thousand causes a page fault
 $= 200 + 7,999,800 \times 1/1000$
 $= 200 + 7999.8 = 8199.8 \text{ ns}$
 $= 8199.8 \times 1/1000 \text{ micro-second [1ms=1000ns]}$
 $= 8.1998 \text{ micro-seconds} \approx 8.2 \text{ micro-seconds}$





Performance of Demand Paging

- EAT= 8.2 micro-seconds
- This is a slowdown by a factor of 40!!
- If want performance degradation < 10 percent
$$220 > 200 + 7999800p \text{ [were EAT= } 200 + 7999800p\text{]}$$
$$220 - 200 > 7999800p$$
$$20 > 7999800p$$
$$p < 20 / 7999800$$
$$p < .0000025$$
- It means one page fault in every 400,000 memory accesses
- **So, we can conclude, it is important to keep the page fault rate very low in a demand paging system**






Basic Page Replacement

1. Find the location of the desired page on disk
2. Find a free frame:
 - If there is a free frame, use it
 - If there is no free frame, use a page replacement algorithm to select a **victim** frame
3. Read the desired page into the (newly) free frame. Update the page and frame tables.
4. Restart the process





Page Replacement






Page Replacement

- Use **modify (dirty) bit** to reduce overhead of page transfers – only modified pages are written to disk
- Page replacement completes separation between logical memory and physical memory – large virtual memory can be provided on a smaller physical memory
- Solve two problems in demand paging implementation:
 - **Frame-allocation algorithm** – how many frames to allocate to each process
 - **Page-replacement algorithm** – select frames to be replaced





Graph of Page Faults Versus The Number of Frames





Page Replacement Algorithms

- Want lowest page-fault rate
- Evaluate algorithm by running it on a particular string of memory references (**reference string**) and computing the number of page faults on that string
- In all our examples, the reference string is

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5





First-In-First-Out (FIFO) Algorithm

- The page which brought first will be replace first
- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- 3 frames (3 pages can be in memory at a time per process)

1	1	4	5	
2	2	1	3	9 page faults
3	3	2	4	

- 4 frames

1	1	5	4	
2	2	1	5	10 page faults
3	3	2		
4	4	3		

- FIFO Replacement – Belady’s Anomaly

- more frames \Rightarrow more page faults





FIFO Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1


7	7	7	2		2	2	4	4	4	0		0	0			7	7	7
	0	0	0		3	3	3	2	2	2		1	1			1	0	0
		1	1		1	0	0	0	3	3		3	2			2	2	1

page frames





FIFO Illustrating Belady's Anomaly





Optimal Algorithm

- Replace page that will not be used for longest period of time
- 4 frames example

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

1	4
2	
3	
4	5


6 page faults

- How do you know this?
- Used for measuring how well your algorithm performs





Optimal Page Replacement





Least Recently Used (LRU) Algorithm

- LRU replaces page that has not been used for the longest time
- Use the recent past to predict the future
- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

1	5
2	
3	5 4
4	3

8 page faults





LRU Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		4	4	4	0			1		1		1		
	0	0	0		0		0	0	3	3			3		0		0		
		1	1		3		3	2	2	2			2		2		7		

page frames





LRU Algorithm (Cont.)

■ Counter implementation

- Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter
- When a page needs to be replaced, look at the counters to determine which has the oldest time-of-access


■ Stack implementation – keep a stack of page numbers in a double link form:

- Page referenced -> move it to the top of stack
 - ▶ bottom of stack will be the LRU page
 - ▶ requires 6 pointers to be changed
- No search for replacement





Use Of A Stack to Record The Most Recent Page References






Thrashing

- It is a situation, if the system spends more time in paging/swapping instead of their execution.
- If a process does not have “enough” frames, the **page-fault rate is very high**. This leads to:
 - low CPU utilization
 - operating system thinks that it needs to increase the degree of multiprogramming
 - another process added to the system
- **Thrashing** \equiv a process is busy swapping pages in and out





Thrashing (Cont.)





Demand Paging and Thrashing


- Why does demand paging work?
- **Locality** model
 - Locality = set of pages in active use
 - Process migrates from one locality to another, e.g. main function, subroutine
 - Localities may overlap
- Why does **thrashing** occur?
 - size of locality > size of allocated frames





Page-Fault Frequency Scheme

- Establish “acceptable” page-fault rate
 - If actual rate too low, process loses frame
 - If actual rate too high, process gains frame






Other Issues -- Prepaging

■ Prepaging

- To reduce the large number of page faults that occurs at process startup
- Prepage all or some of the pages a process will need, before they are referenced
- But if prepaged pages are unused, I/O and memory was wasted
- Assume s pages are prepaged and a fraction α of the s pages is used ($0 \leq \alpha \leq 1$)
 - ▶ Is cost of $s * \alpha$ saved pages faults $>$ or $<$ than the cost of prepaging $s * (1 - \alpha)$ unnecessary pages?
 - ▶ α near zero \Rightarrow prepaging loses
 - ▶ α near one \Rightarrow prepaging wins





Other Issues – Page Size

- **Page size selection** must take into consideration:
 - fragmentation
 - table size
 - I/O overhead
 - locality



End of Chapter 9

